

Jixin (Jerry) Wang

Austin, TX | (469) 556-5559 | jixinwang08@gmail.com | [LinkedIn](#)

EDUCATION

The University of Texas at Austin

B.S. Neuroscience + Computer Science

Austin, TX

Expected May 2028

EXPERIENCE

Moonshot (YC S26) | Software Engineering Intern (Remote)

Jul 2026 – Present

- Owned real-device reliability for the P0 launch-blocking microphone-persistence issue in an ambient-listening iOS app: diagnosed and fixed Swift audio-session recovery races, including a 21-second outage on Bluetooth disconnect that latched into a 300-second permanent refusal, with regression tests for each.
- Designed and built a 19-scenario real-hardware evidence rig in Node.js (lifecycle, interruption, audio route, power, soak) that joins timestamped operator actions to the phone's own ledger and derives outage duration, recovery latency, and per-expectation verdicts; published the first auditable evidence for the blocker across two TestFlight builds, all passing with zero engine restarts.
- Eliminated a quadratic write-ahead-log reconcile that Apple CPU diagnostics showed pinning the app at 82–87% CPU in the background; the WAL now reconciles from disk only after an ambiguous write, cutting two fsyncs and a full replay per commit.
- Fixed release-pipeline defects, including persistent speaker identity silently disabled on every TestFlight build (detected internal distribution via the App Store sandbox receipt instead of the stripped provisioning profile); 23 merged PRs in the first three weeks.

Stèle Inc. | Software Engineering Intern (Remote)

Jul 2026 – Present

- Own the real-time data layer for a neural-sensing wearable: device-to-backend streaming, session management, buffering, and reconnection handling under unreliable network conditions.
- Consolidated 4 duplicated log-classification code paths into a single shared LLM-powered ingest and edit pipeline (structured LLM parsing with keyword fallback), backed by 78 regression tests and shipped through PR review in a shared team codebase.
- Ship production iOS features in SwiftUI/SwiftData, including a voice-driven daily check-in, HealthKit integration, and shared authentication surfaces.

Texas Christian University | Software Engineering Consultant (Fort Worth, TX)

May 2026 – Present

- Sole engineer for a production iOS app deployed at Children's Medical Center Dallas and used daily by physicians, owning architecture, feature development, CI/CD, and production reliability.
- Built the clinical transcription pipeline end to end: live microphone capture via AVFoundation streaming to disk, asynchronous Amazon Transcribe Medical jobs with polling and failure handling, and a custom tokenizer repairing punctuation across variant JSON schemas.
- Implemented AWS Signature Version 4 request signing from scratch in Swift with CryptoKit (key-derivation chain, canonical request construction, payload hashing) for S3 and Transcribe, replacing the AWS SDK with a zero-dependency implementation.
- Integrated Llama 3 8B via a Hugging Face Inference Endpoint for structured health-literacy feedback across four evaluation dimensions, with context-threaded follow-up chat and streaming Markdown rendering.

UNT Vehicle and Edge Computing Lab | Undergraduate Researcher (Denton, TX)

Sep 2024 – May 2026

- Developed the LLM system behind HealthLitPro, an automated audit-and-feedback pipeline assessing clinician fidelity to health-literacy principles; co-authored the resulting paper at IEEE/ACM CHASE 2025.
- Fine-tuned Mistral 7B and Llama 3 8B with LoRA on a curated dataset of 212+ annotated healthcare interactions, and applied RAG to ground model feedback in health-literacy guidelines.

PROJECTS

StrokeSentry | Swift, Apple Vision & Speech frameworks

Lead Author, IEEE ICBAIE 2025

- Built a native iOS app that screens for stroke symptoms via the FAST method entirely on-device, achieving 42.1 ms median inference latency with zero network dependency.
- Designed facial-asymmetry and arm-drift detection over Vision landmarks and body pose (normalized left/right deltas, per-joint confidence gating, 3-second score aggregation) and a speech-clarity model combining recognition confidence, speaking rate, and Jaccard word-set similarity.

SmartNursingHome | JavaScript, TensorFlow.js, Firebase, Swift

President Special Award, DFWIT AI & Startup Competition 2026

- Designed a privacy-first fall detection system where pose estimation runs entirely on-device in the browser and only alarm metadata leaves the patient device; built a three-stage temporal fall detector over torso keypoints and a real-time SwiftUI caregiver dashboard with per-patient Firebase listeners and push notifications.

TECHNICAL SKILLS

Languages: Swift, Python, JavaScript/Node.js, C++, Java

Systems & Cloud: AWS (S3, Transcribe, SigV4), Firebase, real-time streaming, SwiftUI/SwiftData, AVFoundation, CI/CD, TestFlight

ML & AI: PyTorch, Hugging Face (LoRA, RAG), TensorFlow.js, Scikit-learn, LLM integration and evaluation, Apple Vision & Speech